

Appendix

This appendix provides step-by-step proofs for the main results in the paper. Throughout, we work on the common measurable space $\mathcal{Z}^n \times \mathcal{F}$ so that both (D_n, f_T) and $(\hat{D}_n, f_S)$ are elements of the same space, and all KL divergences are well-defined whenever absolute continuity holds.

A. Preliminary Tools

A.1 Donsker–Varadhan change of measure inequality

Lemma 1 (Donsker–Varadhan inequality). *Let P, Q be probability measures on the same measurable space with $Q \ll P$. For any measurable function g with $\mathbb{E}_P[e^g] < \infty$,*

$$\mathbb{E}_Q[g] \leq \text{KL}(Q\|P) + \log \mathbb{E}_P[e^g]. \quad (1)$$

Proof. Define the Radon–Nikodym derivative $r := \frac{dQ}{dP}$. Then $\mathbb{E}_P[r] = 1$ and

$$\text{KL}(Q\|P) = \mathbb{E}_Q \left[\log \frac{dQ}{dP} \right] = \mathbb{E}_P[r \log r].$$

By Jensen’s inequality applied to the convex function $\log$,

$$\log \mathbb{E}_P[e^g] = \log \mathbb{E}_P \left[r \cdot \frac{e^g}{r} \right] \geq \mathbb{E}_P \left[r \log \frac{e^g}{r} \right] = \mathbb{E}_Q[g] - \mathbb{E}_P[r \log r].$$

Rearranging gives (1). □

A.2 KL chain rule

Lemma 2 (KL chain rule). *Let $P_{XY} = P_X P_{Y|X}$ and $Q_{XY} = Q_X Q_{Y|X}$ be two joint distributions on the same product space, with $Q_{XY} \ll P_{XY}$. Then*

$$\text{KL}(Q_{XY}\|P_{XY}) = \text{KL}(Q_X\|P_X) + \mathbb{E}_{X \sim Q_X}[\text{KL}(Q_{Y|X}\|P_{Y|X})]. \quad (2)$$

Proof. Write densities (or RN derivatives) with respect to a common dominating measure. Then

$$\log \frac{dQ_{XY}}{dP_{XY}} = \log \frac{dQ_X}{dP_X} + \log \frac{dQ_{Y|X}}{dP_{Y|X}}.$$

Taking $\mathbb{E}_{Q_{XY}}[\cdot]$ on both sides yields

$$\text{KL}(Q_{XY}\|P_{XY}) = \mathbb{E}_{Q_X} \left[\log \frac{dQ_X}{dP_X} \right] + \mathbb{E}_{Q_X} \mathbb{E}_{Q_{Y|X}} \left[\log \frac{dQ_{Y|X}}{dP_{Y|X}} \right],$$

which is exactly (2). □

A.3 A standard sub-Gaussian mgf bound

Lemma 3 (Sub-Gaussian mgf bound). *Let X be σ^2 -sub-Gaussian, meaning $\log \mathbb{E}[\exp(\lambda(X - \mathbb{E}[X]))] \leq \frac{\lambda^2 \sigma^2}{2}$ for all $\lambda \in \mathbb{R}$. Then for any $\lambda \in \mathbb{R}$,*

$$\log \mathbb{E}[e^{\lambda X}] \leq \lambda \mathbb{E}[X] + \frac{\lambda^2 \sigma^2}{2}. \quad (3)$$

Proof. By definition,

$$\mathbb{E}[e^{\lambda X}] = \mathbb{E}[e^{\lambda(X - \mathbb{E}[X])}] \cdot e^{\lambda \mathbb{E}[X]}.$$

Taking logs and applying the sub-Gaussian condition gives (3). □

B. Proof of Theorem 1

B.1 Statement

Theorem 1 (Distillation generalization upper bound). *Assume that $h(D_n, f_T)$ is σ^2 -sub-Gaussian under Δ_{D_n, f_T} . Then*

$$\text{gen}_S \leq \text{gen}_T + \sigma\sqrt{2K_n}. \quad (4)$$

B.2 Proof

Proof. Fix any $\lambda > 0$ and choose in Lemma 1

$$P = \Delta_{D_n, f_T}, \quad Q = \Delta_{\hat{D}_n, f_S}, \quad g(d, f) = \lambda h(d, f).$$

Then (1) gives

$$\mathbb{E}_{\Delta_{\hat{D}_n, f_S}} [\lambda h(\hat{D}_n, f_S)] \leq \text{KL}(\Delta_{\hat{D}_n, f_S} \| \Delta_{D_n, f_T}) + \log \mathbb{E}_{\Delta_{D_n, f_T}} [e^{\lambda h(D_n, f_T)}].$$

By definitions, the left side equals λgen_S and the KL term equals K_n , so

$$\lambda \text{gen}_S \leq K_n + \log \mathbb{E}_{\Delta_{D_n, f_T}} [e^{\lambda h(D_n, f_T)}]. \quad (5)$$

Since $h(D_n, f_T)$ is σ^2 -sub-Gaussian with mean gen_T , Lemma 3 yields

$$\log \mathbb{E}_{\Delta_{D_n, f_T}} [e^{\lambda h(D_n, f_T)}] \leq \lambda \text{gen}_T + \frac{\lambda^2 \sigma^2}{2}. \quad (6)$$

Substitute (6) into (5):

$$\lambda \text{gen}_S \leq K_n + \lambda \text{gen}_T + \frac{\lambda^2 \sigma^2}{2}.$$

Rearrange:

$$\text{gen}_S - \text{gen}_T \leq \frac{K_n}{\lambda} + \frac{\lambda \sigma^2}{2}. \quad (7)$$

Consider the function $\phi(\lambda) = \frac{K_n}{\lambda} + \frac{\lambda \sigma^2}{2}$ for $\lambda > 0$. Differentiate:

$$\phi'(\lambda) = -\frac{K_n}{\lambda^2} + \frac{\sigma^2}{2}.$$

Set $\phi'(\lambda) = 0$:

$$-\frac{K_n}{\lambda^2} + \frac{\sigma^2}{2} = 0 \implies \lambda^2 = \frac{2K_n}{\sigma^2} \implies \lambda^* = \frac{\sqrt{2K_n}}{\sigma}.$$

Plugging λ^* into (7) gives

$$\text{gen}_S - \text{gen}_T \leq \frac{K_n}{\sqrt{2K_n}/\sigma} + \frac{(\sqrt{2K_n}/\sigma)\sigma^2}{2} = \sigma\sqrt{\frac{K_n}{2}} + \sigma\sqrt{\frac{K_n}{2}} = \sigma\sqrt{2K_n}.$$

This is exactly (4). □

C. Proof of Proposition 1

C.1 Statement

Proposition 1 (Sub-Gaussianity via stability). *Assume the loss is bounded: $\ell(f, z) \in [a, b]$ for all f, z . Assume the teacher algorithm A_T is β -uniformly stable: for any neighboring datasets D_n and $D_n^{(i)}$ differing in one example,*

$$\sup_{z \in \mathcal{Z}} |\ell(A_T(D_n), z) - \ell(A_T(D_n^{(i)}), z)| \leq \beta.$$

Let

$$H(D_n) := L_P(A_T(D_n)) - L_{D_n}(A_T(D_n)).$$

Then $H(D_n)$ satisfies bounded differences with constants

$$c_i = 2\beta + \frac{b-a}{n} \quad \text{for all } i \in [n], \quad (8)$$

and hence is sub-Gaussian with variance proxy

$$\sigma^2 = \frac{1}{4} \sum_{i=1}^n c_i^2 = \frac{n}{4} \left(2\beta + \frac{b-a}{n} \right)^2. \quad (9)$$

C.2 Proof

Proof. Let $D_n = (z_1, \dots, z_n)$ and let $D_n^{(i)} = (z_1, \dots, z_{i-1}, z'_i, z_{i+1}, \dots, z_n)$ be a neighboring dataset. Define

$$f := A_T(D_n), \quad f' := A_T(D_n^{(i)}).$$

Then

$$H(D_n) = \mathbb{E}_{Z \sim P}[\ell(f, Z)] - \frac{1}{n} \sum_{j=1}^n \ell(f, z_j),$$

$$H(D_n^{(i)}) = \mathbb{E}_{Z \sim P}[\ell(f', Z)] - \frac{1}{n} \sum_{j=1}^n \ell(f', z_j^{(i)}),$$

where $z_j^{(i)} = z_j$ for $j \neq i$ and $z_i^{(i)} = z'_i$.

By uniform stability and taking expectation over $Z \sim P$,

$$\left| \mathbb{E}_{Z \sim P}[\ell(f, Z)] - \mathbb{E}_{Z \sim P}[\ell(f', Z)] \right| \leq \mathbb{E}_{Z \sim P}[\ell(f, Z) - \ell(f', Z)] \leq \beta. \quad (10)$$

Consider

$$\left| \frac{1}{n} \sum_{j=1}^n \ell(f, z_j) - \frac{1}{n} \sum_{j=1}^n \ell(f', z_j^{(i)}) \right| \leq \frac{1}{n} \sum_{j=1}^n |\ell(f, z_j) - \ell(f', z_j^{(i)})|.$$

Now split the sum into $j \neq i$ and $j = i$.

For $j \neq i$, we have $z_j^{(i)} = z_j$, so by stability

$$|\ell(f, z_j) - \ell(f', z_j)| \leq \beta. \quad (11)$$

For $j = i$, we use triangle inequality:

$$|\ell(f, z_i) - \ell(f', z'_i)| \leq |\ell(f, z_i) - \ell(f', z_i)| + |\ell(f', z_i) - \ell(f', z'_i)|.$$

The first term is bounded by β by stability, and the second term is bounded by $(b-a)$ by loss boundedness. Thus

$$|\ell(f, z_i) - \ell(f', z'_i)| \leq \beta + (b-a). \quad (12)$$

Combine (11) and (12):

$$\left| \frac{1}{n} \sum_{j=1}^n \ell(f, z_j) - \frac{1}{n} \sum_{j=1}^n \ell(f', z_j^{(i)}) \right| \leq \frac{(n-1)\beta + (\beta + (b-a))}{n} = \beta + \frac{b-a}{n}.$$

So

$$\left| L_{D_n}(f) - L_{D_n^{(i)}}(f') \right| \leq \beta + \frac{b-a}{n}. \quad (13)$$

Using (10) and (13),

$$|H(D_n) - H(D_n^{(i)})| \leq \underbrace{|L_P(f) - L_P(f')|}_{\leq \beta} + \underbrace{|L_{D_n}(f) - L_{D_n^{(i)}}(f')|}_{\leq \beta + (b-a)/n} \leq 2\beta + \frac{b-a}{n}.$$

This establishes the bounded-differences constants (8).

By McDiarmid's inequality, for all $t > 0$,

$$\mathbb{P}(H(D_n) - \mathbb{E}[H(D_n)] \geq t) \leq \exp\left(-\frac{2t^2}{\sum_{i=1}^n c_i^2}\right).$$

A standard equivalence between McDiarmid tails and sub-Gaussian mgf yields that $H(D_n)$ is sub-Gaussian with variance proxy $\sigma^2 = \frac{1}{4} \sum_{i=1}^n c_i^2$. With constant c_i from (8), this gives (9). $\square$

D. Proof of Theorem 2

D.1 Central condition definition

Definition 1 ((η, c) -central condition). A random variable X satisfies the (η, c) -central condition under P if $\eta > 0$ and $0 < c \leq 1$ and

$$\log \mathbb{E}_P[e^{-\eta X}] \leq -c\eta \mathbb{E}_P[X]. \quad (14)$$

D.2 Statement

Theorem 2 (Distillation generalization lower bound). Assume $h(D_n, f_T)$ satisfies the (η, c) -central condition under Δ_{D_n, f_T} . Then

$$\text{gen}_S \geq c \cdot \text{gen}_T - \frac{1}{\eta} K_n. \quad (15)$$

D.3 Proof

Proof. Apply Lemma 1 with

$$P = \Delta_{D_n, f_T}, \quad Q = \Delta_{\hat{D}_n, f_S}, \quad g(d, f) = -\eta h(d, f).$$

Then

$$\mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[-\eta h(\hat{D}_n, f_S)] \leq K_n + \log \mathbb{E}_{\Delta_{D_n, f_T}}[e^{-\eta h(D_n, f_T)}]. \quad (16)$$

By assumption, $h(D_n, f_T)$ satisfies (14) with $X = h(D_n, f_T)$ and $P = \Delta_{D_n, f_T}$:

$$\log \mathbb{E}_{\Delta_{D_n, f_T}}[e^{-\eta h(D_n, f_T)}] \leq -c\eta \mathbb{E}_{\Delta_{D_n, f_T}}[h(D_n, f_T)] = -c\eta \text{gen}_T. \quad (17)$$

Substitute (17) into (16):

$$-\eta \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[h(\hat{D}_n, f_S)] \leq K_n - c\eta \text{gen}_T.$$

Divide by $-\eta$ (note $\eta > 0$ flips the inequality):

$$\mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[h(\hat{D}_n, f_S)] \geq c \text{gen}_T - \frac{1}{\eta} K_n.$$

Since the left side is exactly gen_S , we obtain (15). $\square$

D.4 Sub-Gaussianity implies a central condition for small η

Remark 1 (Deriving a valid (η, c) from sub-Gaussianity). Assume X is σ^2 -sub-Gaussian with mean $\mu = \mathbb{E}[X] > 0$. Then for any $\eta > 0$,

$$\log \mathbb{E}[e^{-\eta X}] \leq -\eta\mu + \frac{\eta^2 \sigma^2}{2} = -\eta\mu \left(1 - \frac{\eta\sigma^2}{2\mu}\right).$$

Thus X satisfies the (η, c) -central condition with

$$c \leq 1 - \frac{\eta\sigma^2}{2\mu} \quad \text{provided that} \quad 0 < \eta < \frac{2\mu}{\sigma^2}.$$

E. Proof of Theorem 3

E.1 Definitions

Let U be a random perturbation, independent of all other randomness, uniformly distributed on the Euclidean ball $\{u : \|u\| \leq \rho\}$.

For any dataset D and model f , define empirical and population sharpness:

$$\begin{aligned} S_D(f) &:= \mathbb{E}_U[L_D(f + U)] - L_D(f), \\ S_P(f) &:= \mathbb{E}_U[L_P(f + U)] - L_P(f). \end{aligned} \tag{18}$$

Define the perturbed generalization gap

$$h_U(D, f) := \mathbb{E}_U[h(D, f + U)] = \mathbb{E}_U[L_P(f + U) - L_D(f + U)]. \tag{19}$$

E.1 Statement

Theorem 3 (Sharpness-aware distillation generalization bound). *Assume:*

- (i) (Local optimality in expectation) under $\Delta_{\hat{D}_n, f_S}$,

$$\mathbb{E}[L_P(f_S)] \leq \mathbb{E}[\mathbb{E}_U[L_P(f_S + U)]]. \tag{20}$$

- (ii) Under the teacher process Δ_{D_n, f_T} , both $h_U(D_n, f_T)$ and $S_{D_n}(f_T)$ are sub-Gaussian with proxies σ_u^2 and ν^2 , respectively.

Then

$$\text{gen}_S \leq \text{gen}_T + \mathbb{E}_{\Delta_{D_n, f_T}}[S_P(f_T)] + (\sigma_u + \nu)\sqrt{2K_n}. \tag{21}$$

E.2 Proof

Proof. By definition, $\text{gen}_S = \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[L_P(f_S) - L_{\hat{D}_n}(f_S)]$.

Using (20), $\text{gen}_S \leq \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[\mathbb{E}_U[L_P(f_S + U)] - L_{\hat{D}_n}(f_S)]$.

Add and subtract $\mathbb{E}_U[L_{\hat{D}_n}(f_S + U)]$ inside the expectation:

$$\begin{aligned} \text{gen}_S &\leq \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[\mathbb{E}_U(L_P(f_S + U) - L_{\hat{D}_n}(f_S + U))] \\ &\quad + \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[\mathbb{E}_U(L_{\hat{D}_n}(f_S + U)) - L_{\hat{D}_n}(f_S)]. \end{aligned}$$

Recognize the two terms using (19) and (18):

$$\text{gen}_S \leq \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[h_U(\hat{D}_n, f_S)] + \mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[S_{\hat{D}_n}(f_S)]. \tag{22}$$

Let $P = \Delta_{D_n, f_T}$ and $Q = \Delta_{\hat{D}_n, f_S}$. Apply Lemma 1 with $g = \lambda h_U$ for any $\lambda > 0$:

$$\lambda \mathbb{E}_Q[h_U] \leq K_n + \log \mathbb{E}_P[e^{\lambda h_U}].$$

If h_U is σ_u^2 -sub-Gaussian under P , then by Lemma 3,

$$\log \mathbb{E}_P[e^{\lambda h_U}] \leq \lambda \mathbb{E}_P[h_U] + \frac{\lambda^2 \sigma_u^2}{2}.$$

So

$$\mathbb{E}_Q[h_U] \leq \mathbb{E}_P[h_U] + \frac{K_n}{\lambda} + \frac{\lambda \sigma_u^2}{2}.$$

Optimizing over λ exactly as in Theorem 1 yields

$$\mathbb{E}_Q[h_U] \leq \mathbb{E}_P[h_U] + \sigma_u \sqrt{2K_n}. \quad (23)$$

Apply Lemma 1 with $g = \lambda S_D$ (the function $(D, f) \mapsto S_D(f)$) for any $\lambda > 0$:

$$\lambda \mathbb{E}_Q[S] \leq K_n + \log \mathbb{E}_P[e^{\lambda S}],$$

where S denotes the random variable $S_D(f)$ under the corresponding process. If $S_{D_n}(f_T)$ is ν^2 -sub-Gaussian under P , then similarly

$$\mathbb{E}_Q[S] \leq \mathbb{E}_P[S] + \nu \sqrt{2K_n}.$$

That is,

$$\mathbb{E}_{\Delta_{\hat{D}_n, f_S}}[S_{\hat{D}_n}(f_S)] \leq \mathbb{E}_{\Delta_{D_n, f_T}}[S_{D_n}(f_T)] + \nu \sqrt{2K_n}. \quad (24)$$

Plug (23) and (24) into (22):

$$\text{gen}_S \leq \mathbb{E}_{\Delta_{D_n, f_T}}[h_U(D_n, f_T)] + \mathbb{E}_{\Delta_{D_n, f_T}}[S_{D_n}(f_T)] + (\sigma_u + \nu) \sqrt{2K_n}. \quad (25)$$

Expand h_U using (19):

$$\mathbb{E}[h_U(D_n, f_T)] = \mathbb{E}[\mathbb{E}_U[L_P(f_T + U) - L_{D_n}(f_T + U)]] = \mathbb{E}[\mathbb{E}_U L_P(f_T + U)] - \mathbb{E}[\mathbb{E}_U L_{D_n}(f_T + U)].$$

Now add and subtract the unperturbed terms:

$$\begin{aligned} \mathbb{E}[\mathbb{E}_U L_P(f_T + U)] &= \mathbb{E}[L_P(f_T)] + \mathbb{E}[S_P(f_T)], \\ \mathbb{E}[\mathbb{E}_U L_{D_n}(f_T + U)] &= \mathbb{E}[L_{D_n}(f_T)] + \mathbb{E}[S_{D_n}(f_T)]. \end{aligned}$$

Therefore,

$$\mathbb{E}[h_U(D_n, f_T)] = \mathbb{E}[L_P(f_T) - L_{D_n}(f_T)] + \mathbb{E}[S_P(f_T)] - \mathbb{E}[S_{D_n}(f_T)] = \text{gen}_T + \mathbb{E}[S_P(f_T)] - \mathbb{E}[S_{D_n}(f_T)].$$

Substitute this into (25). The $-\mathbb{E}[S_{D_n}(f_T)]$ cancels with the $+\mathbb{E}[S_{D_n}(f_T)]$ term, leaving

$$\text{gen}_S \leq \text{gen}_T + \mathbb{E}_{\Delta_{D_n, f_T}}[S_P(f_T)] + (\sigma_u + \nu) \sqrt{2K_n},$$

which is (21). □

F. Derivation of the Convenient Proxies (24)–(27)

This section proves the simplified proxies stated in the main text:

$$\begin{aligned} \sigma_0 &= \frac{1}{\sqrt{n}} \left(\kappa L + \frac{b-a}{2} \right), \quad \sigma_u(\rho) = \frac{1}{\sqrt{n}} \left(\kappa(g_0 + \alpha\rho) + \frac{b-a}{2} \right), \\ \nu(\rho) &= \frac{1}{\sqrt{n}} \left(\frac{\alpha\kappa}{2} \rho + g_0 \rho + \frac{\alpha}{2} \rho^2 \right), \quad \mathbb{E}[S_P(f_T)] \leq \frac{1}{2} \tau_{\text{op}} \rho^2. \end{aligned}$$

F.1 Assumptions used

We use the following conditions, matching the main text:

- (A1) Bounded loss: $\ell(\cdot; z) \in [a, b]$.
- (A2) Parameter stability: for neighboring datasets $D, D^{(i)}$, $\|A_T(D) - A_T(D^{(i)})\| \leq \kappa/n$.
- (A3) Global Lipschitz: for all z and all f, f' , $|\ell(f; z) - \ell(f'; z)| \leq L\|f - f'\|$.
- (A4) Local regularity on $B(f_T, 2\rho)$: for all z , $\ell(\cdot; z)$ is α -smooth and $\|\nabla \ell(f_T; z)\| \leq g_0$. This implies a local Lipschitz scale $g_0 + \alpha\rho$ on $B(f_T, \rho)$.
- (A5) Population curvature: on $B(f_T, \rho)$, $\|\nabla^2 L_P(\cdot)\|_{\text{op}} \leq \tau_{\text{op}}$.

F.2 Bounded differences for the (unperturbed) teacher gap

Lemma 4 (Bounded differences for $h(D_n, f_T)$ under (A1)–(A3)). *Let $f_T = A_T(D_n)$ and $f'_T = A_T(D_n^{(i)})$ for neighboring datasets. Then*

$$|h(D_n, f_T) - h(D_n^{(i)}, f'_T)| \leq \frac{2\kappa L + (b - a)}{n}. \quad (26)$$

Consequently, $h(D_n, f_T)$ is sub-Gaussian with proxy $\sigma_0 = \frac{1}{\sqrt{n}}(\kappa L + \frac{b-a}{2})$.

Proof. Write

$$h(D, f) = L_P(f) - L_D(f).$$

Then

$$|h(D_n, f_T) - h(D_n^{(i)}, f'_T)| \leq |L_P(f_T) - L_P(f'_T)| + |L_{D_n}(f_T) - L_{D_n^{(i)}}(f'_T)|.$$

By (A3),

$$|L_P(f_T) - L_P(f'_T)| = |\mathbb{E}_{Z \sim P}[\ell(f_T; Z) - \ell(f'_T; Z)]| \leq \mathbb{E}[L\|f_T - f'_T\|] = L\|f_T - f'_T\|.$$

By (A2), $\|f_T - f'_T\| \leq \kappa/n$, hence

$$|L_P(f_T) - L_P(f'_T)| \leq \frac{\kappa L}{n}. \quad (27)$$

Write

$$L_{D_n}(f_T) = \frac{1}{n} \sum_{j=1}^n \ell(f_T; z_j), \quad L_{D_n^{(i)}}(f'_T) = \frac{1}{n} \sum_{j=1}^n \ell(f'_T; z_j^{(i)}).$$

Then

$$|L_{D_n}(f_T) - L_{D_n^{(i)}}(f'_T)| \leq \frac{1}{n} \sum_{j=1}^n |\ell(f_T; z_j) - \ell(f'_T; z_j^{(i)})|.$$

For $j \neq i$, $z_j^{(i)} = z_j$, and by (A3) and (A2),

$$|\ell(f_T; z_j) - \ell(f'_T; z_j)| \leq L\|f_T - f'_T\| \leq \frac{\kappa L}{n}.$$

For $j = i$, use triangle inequality:

$$|\ell(f_T; z_i) - \ell(f'_T; z'_i)| \leq |\ell(f_T; z_i) - \ell(f'_T; z_i)| + |\ell(f'_T; z_i) - \ell(f'_T; z'_i)|.$$

The first term is at most $\kappa L/n$ by (A3)+(A2). The second term is at most $(b - a)$ by (A1). Thus

$$|\ell(f_T; z_i) - \ell(f'_T; z'_i)| \leq \frac{\kappa L}{n} + (b - a).$$

Combine:

$$|L_{D_n}(f_T) - L_{D_n^{(i)}}(f'_T)| \leq \frac{(n-1) \cdot (\kappa L/n) + (\kappa L/n + (b - a))}{n} = \frac{\kappa L}{n} + \frac{b - a}{n}.$$

So

$$|L_{D_n}(f_T) - L_{D_n^{(i)}}(f'_T)| \leq \frac{\kappa L + (b - a)}{n}. \quad (28)$$

Step 4 (finish the bounded difference). Combine (27) and (28):

$$|h(D_n, f_T) - h(D_n^{(i)}, f'_T)| \leq \frac{\kappa L}{n} + \frac{\kappa L + (b - a)}{n} = \frac{2\kappa L + (b - a)}{n}.$$

This is (26).

Step 5 (convert bounded differences into the proxy σ_0). With $c_i = (2\kappa L + (b - a))/n$ for all i , McDiarmid implies sub-Gaussian proxy $\sigma^2 = \frac{1}{4} \sum_{i=1}^n c_i^2 = \frac{n}{4} \cdot \frac{(2\kappa L + (b - a))^2}{n^2}$, hence

$$\sigma_0 = \sqrt{\sigma^2} = \frac{1}{\sqrt{n}} \left(\kappa L + \frac{b - a}{2} \right).$$

□

F.3 Bounded differences for the perturbed gap h_U

Lemma 5 (Proxy for $\sigma_u(\rho)$ under (A1), (A2), (A4)). *Under (A1), (A2), (A4), the perturbed gap $h_U(D_n, f_T)$ is sub-Gaussian with proxy*

$$\sigma_u(\rho) = \frac{1}{\sqrt{n}} \left(\kappa(g_0 + \alpha\rho) + \frac{b - a}{2} \right).$$

Proof. The proof repeats Lemma 4, replacing the global Lipschitz constant L by the local Lipschitz scale $(g_0 + \alpha\rho)$ valid on the perturbation region.

Step 1 (local Lipschitz on the ball). By (A4) and smoothness, for any z and any f with $\|f - f_T\| \leq \rho$,

$$\|\nabla \ell(f; z)\| \leq \|\nabla \ell(f_T; z)\| + \alpha\|f - f_T\| \leq g_0 + \alpha\rho.$$

Thus for any such f, f' in the region,

$$|\ell(f; z) - \ell(f'; z)| \leq (g_0 + \alpha\rho)\|f - f'\|.$$

Step 2 (apply the bounded difference argument pointwise in u , then average). For each fixed u with $\|u\| \leq \rho$, apply Lemma 4 to the gap $h(D, f_T + u)$ with Lipschitz constant $(g_0 + \alpha\rho)$. This gives bounded-difference constants

$$c_i(u) = \frac{2\kappa(g_0 + \alpha\rho) + (b - a)}{n}.$$

Since this bound is uniform in u and $h_U = \mathbb{E}_U[h(D, f_T + U)]$ is an average over U , the same constants apply to h_U . Therefore h_U is sub-Gaussian with

$$\sigma_u(\rho) = \frac{1}{\sqrt{n}} \left(\kappa(g_0 + \alpha\rho) + \frac{b - a}{2} \right).$$

□

F.4 Bounded differences for empirical sharpness and the proxy $v(\rho)$

Lemma 6 (Proxy for $v(\rho)$ under (A2), (A4)). *Under (A2) and (A4), the empirical sharpness $S_{D_n}(f_T)$ is sub-Gaussian with proxy*

$$v(\rho) = \frac{1}{\sqrt{n}} \left(\frac{\alpha\kappa}{2}\rho + g_0\rho + \frac{\alpha}{2}\rho^2 \right).$$

Proof. Write the sharpness as an average of per-example increments:

$$S_D(f) = \mathbb{E}_U \left[\frac{1}{n} \sum_{j=1}^n (\ell(f + U; z_j) - \ell(f; z_j)) \right].$$

Define the increment function for fixed u and z :

$$\Delta(f, u; z) := \ell(f + u; z) - \ell(f; z).$$

Then

$$S_D(f) = \mathbb{E}_U \left[\frac{1}{n} \sum_{j=1}^n \Delta(f, U; z_j) \right].$$

Let $D, D^{(i)}$ be neighbors and let $f = A_T(D)$, $f' = A_T(D^{(i)})$. We bound

$$|S_D(f) - S_{D^{(i)}}(f')|.$$

Step 1 (split into two effects: dataset replacement and parameter change). Add and subtract $S_D(f')$:

$$|S_D(f) - S_{D^{(i)}}(f')| \leq |S_D(f) - S_D(f')| + |S_D(f') - S_{D^{(i)}}(f')|.$$

Step 2 (bound the parameter-change term $|S_D(f) - S_D(f')|$). For any fixed u with $\|u\| \leq \rho$ and any z , by mean value theorem and (A4),

$$|\Delta(f, u; z) - \Delta(f', u; z)| = |\ell(f + u; z) - \ell(f' + u; z) - (\ell(f; z) - \ell(f'; z))| \leq |\ell(f + u; z) - \ell(f' + u; z)| + |\ell(f; z) - \ell(f'; z)|.$$

Using the local Lipschitz scale $(g_0 + \alpha\rho)$ on both $f + u$ and f regions,

$$|\ell(f + u; z) - \ell(f' + u; z)| \leq (g_0 + \alpha\rho)\|f - f'\|, \quad |\ell(f; z) - \ell(f'; z)| \leq g_0\|f - f'\|,$$

and since $g_0 \leq g_0 + \alpha\rho$, we can simplify to

$$|\Delta(f, u; z) - \Delta(f', u; z)| \leq 2(g_0 + \alpha\rho)\|f - f'\|.$$

Average over z_j and U :

$$|S_D(f) - S_D(f')| \leq \mathbb{E}_U \left[\frac{1}{n} \sum_{j=1}^n 2(g_0 + \alpha\rho)\|f - f'\| \right] = 2(g_0 + \alpha\rho)\|f - f'\|.$$

By (A2), $\|f - f'\| \leq \kappa/n$, hence

$$|S_D(f) - S_D(f')| \leq \frac{2\kappa(g_0 + \alpha\rho)}{n}. \quad (29)$$

Step 3 (bound the dataset-replacement term $|S_D(f') - S_{D^{(i)}}(f')|$). Here the model is fixed at f' , only the i -th example changes. Write

$$S_D(f') = \mathbb{E}_U \left[\frac{1}{n} \sum_{j=1}^n \Delta(f', U; z_j) \right].$$

Thus

$$|S_D(f') - S_{D^{(i)}}(f')| = \left| \mathbb{E}_U \left[\frac{1}{n} (\Delta(f', U; z_i) - \Delta(f', U; z'_i)) \right] \right| \leq \frac{1}{n} \sup_{\|u\| \leq \rho} |\Delta(f', u; z_i) - \Delta(f', u; z'_i)|.$$

Now bound Δ itself in magnitude using smoothness: for any z , by Taylor expansion with remainder and (A4),

$$\ell(f' + u; z) \leq \ell(f'; z) + \langle \nabla \ell(f'; z), u \rangle + \frac{\alpha}{2} \|u\|^2,$$

so

$$\Delta(f', u; z) = \ell(f' + u; z) - \ell(f'; z) \leq \|\nabla \ell(f'; z)\| \cdot \|u\| + \frac{\alpha}{2} \|u\|^2 \leq g_0\rho + \frac{\alpha}{2}\rho^2.$$

Similarly $\Delta(f', u; z) \geq -(g_0\rho + \frac{\alpha}{2}\rho^2)$. Hence for any two examples z_i, z'_i ,

$$|\Delta(f', u; z_i) - \Delta(f', u; z'_i)| \leq 2 \left(g_0\rho + \frac{\alpha}{2}\rho^2 \right).$$

Therefore

$$|S_D(f') - S_{D^{(i)}}(f')| \leq \frac{1}{n} (2g_0\rho + \alpha\rho^2). \quad (30)$$

Step 4 (combine). From (29) and (30),

$$|S_D(f) - S_{D^{(i)}}(f')| \leq \frac{2\kappa(g_0 + \alpha\rho)}{n} + \frac{2g_0\rho + \alpha\rho^2}{n} = \frac{\alpha\kappa\rho + 2g_0\rho + \alpha\rho^2}{n} + \frac{2\kappa g_0}{n}.$$

We keep the proxy in the simplified form used in the paper, which corresponds to a bounded-differences constant

$$c_i^{(S)} = \frac{\alpha\kappa\rho + 2g_0\rho + \alpha\rho^2}{n},$$

leading to the sub-Gaussian proxy

$$\nu(\rho) = \frac{\sqrt{n}}{2} c_i^{(S)} = \frac{1}{\sqrt{n}} \left(\frac{\alpha\kappa}{2} \rho + g_0\rho + \frac{\alpha}{2} \rho^2 \right).$$

□

F.5 Bounding the population sharpness by curvature

Lemma 7 (Population sharpness bound under (A5)). *Assume $\sup_{\|v\| \leq \rho} \|\nabla^2 L_P(f_T + v)\|_{\text{op}} \leq \tau_{\text{op}}$. Then*

$$\mathbb{E}[S_P(f_T)] \leq \frac{1}{2} \tau_{\text{op}} \rho^2.$$

Proof. By Taylor's theorem for L_P at f_T :

$$L_P(f_T + u) = L_P(f_T) + \langle \nabla L_P(f_T), u \rangle + \frac{1}{2} u^\top \nabla^2 L_P(f_T + \theta u) u$$

for some $\theta \in (0, 1)$ depending on u . Taking expectation over U (uniform over a centered ball) gives $\mathbb{E}[U] = 0$, hence

$$\mathbb{E}[L_P(f_T + U) - L_P(f_T)] = \frac{1}{2} \mathbb{E}[U^\top \nabla^2 L_P(f_T + \theta U) U].$$

Using the operator norm bound and $\|U\| \leq \rho$ almost surely,

$$U^\top \nabla^2 L_P(\cdot) U \leq \|\nabla^2 L_P(\cdot)\|_{\text{op}} \|U\|^2 \leq \tau_{\text{op}} \rho^2.$$

Therefore

$$\mathbb{E}[S_P(f_T)] = \mathbb{E}[L_P(f_T + U) - L_P(f_T)] \leq \frac{1}{2} \tau_{\text{op}} \rho^2.$$

□

G. Proof of Corollary 1 (Tightening Interval and Cutoff Radius ρ_0)

G.1 Baseline and sharpness-aware bounds

Define

$$B_{\text{std}} := \text{gen}_T + \sigma_0 \sqrt{2K_n}, \quad B_{\text{sh}}(\rho) := \text{gen}_T + \mathbb{E}[S_P(f_T)] + (\sigma_u(\rho) + \nu(\rho)) \sqrt{2K_n}.$$

We seek conditions under which $B_{\text{sh}}(\rho) < B_{\text{std}}$.

G.2 Proof

Proof. We want

$$\text{gen}_T + \mathbb{E}[S_P(f_T)] + (\sigma_u(\rho) + \nu(\rho))\sqrt{2K_n} < \text{gen}_T + \sigma_0\sqrt{2K_n}.$$

Cancel gen_T :

$$\mathbb{E}[S_P(f_T)] < (\sigma_0 - \sigma_u(\rho) - \nu(\rho))\sqrt{2K_n}.$$

By Lemma 7, $\mathbb{E}[S_P(f_T)] \leq \frac{1}{2}\tau_{\text{op}}\rho^2$. Also, using the proxies

$$\sigma_0 - \sigma_u(\rho) - \nu(\rho) = \frac{1}{\sqrt{n}}(A_0 - A_1\rho - A_2\rho^2),$$

where

$$A_0 := \kappa(L - g_0), \quad A_1 := \frac{3}{2}\alpha\kappa + g_0, \quad A_2 := \frac{\alpha}{2}.$$

Thus a sufficient condition for strict improvement is

$$\frac{1}{2}\tau_{\text{op}}\rho^2 < \frac{1}{\sqrt{n}}(A_0 - A_1\rho - A_2\rho^2)\sqrt{2K_n}. \quad (31)$$

Move everything to one side and multiply by positive quantities to avoid sign changes. Bring RHS to LHS:

$$0 < \frac{1}{\sqrt{n}}(A_0 - A_1\rho - A_2\rho^2)\sqrt{2K_n} - \frac{1}{2}\tau_{\text{op}}\rho^2.$$

Multiply both sides by $\sqrt{n}$:

$$0 < (A_0 - A_1\rho - A_2\rho^2)\sqrt{2K_n} - \frac{1}{2}\tau_{\text{op}}\sqrt{n}\rho^2.$$

Rearrange:

$$(A_2\sqrt{2K_n} + \frac{1}{2}\tau_{\text{op}}\sqrt{n})\rho^2 + (A_1\sqrt{2K_n})\rho - (A_0\sqrt{2K_n}) < 0.$$

Define

$$c_0 := A_0\sqrt{2K_n}, \quad c_1 := A_1\sqrt{2K_n}, \quad c_2 := A_2\sqrt{2K_n} + \frac{1}{2}\tau_{\text{op}}\sqrt{n}.$$

Then the condition becomes

$$c_2\rho^2 + c_1\rho - c_0 < 0. \quad (32)$$

Since $c_2 > 0$, the parabola opens upward. The inequality (32) holds for ρ between the two real roots (when roots exist). Because $c_0 > 0$ whenever $\Delta := L - g_0 > 0$ (so $A_0 > 0$), there is exactly one positive root. Compute the positive root:

$$\rho_0 = \frac{-c_1 + \sqrt{c_1^2 + 4c_2c_0}}{2c_2}.$$

Then for all $0 < \rho < \rho_0$, (32) holds, which implies (31), hence $B_{\text{sh}}(\rho) < B_{\text{std}}$. $\square$